

COMP498G/691G COMPUTER VISION

LECTURE 4 EDGE DETECTION



Today's Lecture

- Edge detection
 - Slides acknowledgment: A. Farhadi, S. Seitz
- Questions

Edges and Scale


SHADOW

From Sandlot Science

Today's reading

- Cipolla & Gee on edge detection (on class website)
- Szeliski Ch. 3.2, 3.4, 3.5, 4.2

Origin of Edges




Edges are caused by a variety of factors

Characterizing edges

An edge is a place of rapid change in the image intensity function


image



Characterizing edges

An edge is a place of rapid change in the image intensity function

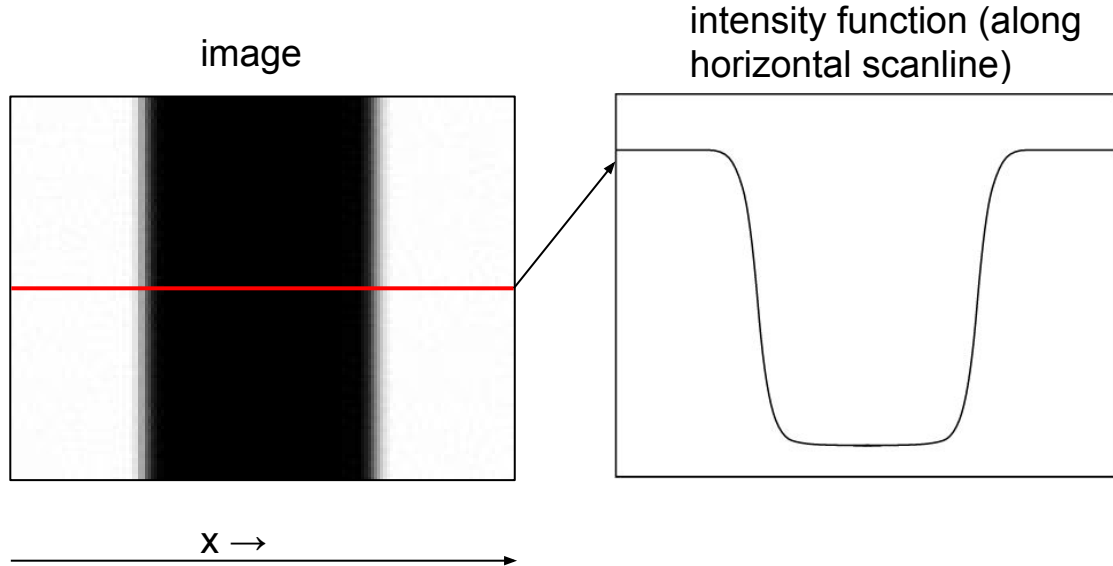
image



$x \rightarrow$

Characterizing edges


An edge is a place of rapid change in the image intensity function



Characterizing edges


An edge is a place of rapid change in the image intensity function

image




$x \rightarrow$

intensity function (along horizontal scanline)




first derivative



Characterizing edges


An edge is a place of rapid change in the image intensity function

image




$x \rightarrow$

intensity function (along horizontal scanline)



first derivative



edges correspond to
extrema of derivative

Image Gradient



$$\frac{\partial f}{\partial x}$$





Image Gradient



$$\frac{\partial f}{\partial x}$$



- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

Image Gradient




$$\frac{\partial f}{\partial x}$$




- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of rapid change in intensity

Image Gradient




- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of rapid change in intensity




$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

Image Gradient



- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of rapid change in intensity




$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$




$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


Image Gradient




- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of rapid change in intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$




$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$




$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Image Gradient



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$


Image Gradient



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient direction is given by: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$


Image Gradient



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient direction is given by: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$ How does this relate to the direction of the edge?

Image Gradient



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient direction is given by: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$ How does this relate to the direction of the edge?
- The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The discrete gradient

How can we differentiate a *digital* image $F[x,y]$?

The discrete gradient

How can we differentiate a digital image $F[x,y]$?

- Option 1: reconstruct a continuous image, then take gradient

The discrete gradient

How can we differentiate a digital image $F[x,y]$?

- Option 1: reconstruct a continuous image, then take gradient
- Option 2: take discrete derivative (“finite difference”)

The discrete gradient

How can we differentiate a digital image $F[x,y]$?

- Option 1: reconstruct a continuous image, then take gradient
- Option 2: take discrete derivative (“finite difference”)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$



Step size is 1


The discrete gradient

How can we differentiate a digital image $F[x,y]$?

- Option 1: reconstruct a continuous image, then take gradient
- Option 2: take discrete derivative (“finite difference”)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a cross-correlation?



H

The Sobel operator

Better approximations of the derivatives exist

- The Sobel operators below are very commonly used

$$\frac{1}{8}$$

-1	0	1
-2	0	2
-1	0	1

s_x

The Sobel operator

Better approximations of the derivatives exist

- The Sobel operators below are very commonly used

$$\frac{1}{8}$$

-1	0	1
-2	0	2
-1	0	1

s_x

$$\frac{1}{8}$$

1	2	1
0	0	0
-1	-2	-1

s_y

The Sobel operator

Better approximations of the derivatives exist

- The Sobel operators below are very commonly used

$$\frac{1}{8}$$

-1	0	1
-2	0	2
-1	0	1

s_x

$$\frac{1}{8}$$

1	2	1
0	0	0
-1	-2	-1

s_y

- The standard definition of the Sobel operator omits the $1/8$ term
 - doesn't make a difference for edge detection
 - the $1/8$ term is needed to get the right gradient value, however


The Sobel operator



Original



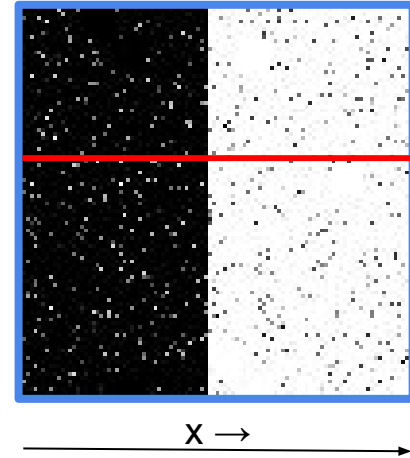
Magnitude



Orientation

Effects of noise



Consider a single row or column of the image



Effects of noise

Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal




Effects of noise



Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal

$$f(x)$$






$$\frac{d}{dx} f(x)$$



Effects of noise

Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal




Where is the edge?


Effects of noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbours
 - Generally, the larger the noise the stronger the response
- What can we do about it?


Solution: smooth first



Solution: smooth first




Solution: smooth first




Solution: smooth first

Sigma = 50


f




h




$h \star f$



$\frac{\partial}{\partial x}(h \star f)$




Solution: smooth first



Where is the edge?

Solution: smooth first




Where is the edge? Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution

- Differentiation is convolution, and convolution is associative: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$


Derivative theorem of convolution

- Differentiation is convolution, and convolution is associative: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$
- This saves us one operation:




Derivative theorem of convolution



- Differentiation is convolution, and convolution is associative: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$
- This saves us one operation:




Remember: Derivative of Gaussian filter



x-direction




y-direction




Laplacian of Gaussian

Consider $\frac{\partial^2}{\partial x^2}(h \star f)$


f



$\frac{\partial^2}{\partial x^2}h$




$(\frac{\partial^2}{\partial x^2}h) \star f$




Where is the edge? Zero-crossings of bottom graph

2D edge detection filters



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_{\sigma}(u, v)$$

∇^2 is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

filter demo

Edge detection by subtraction



original

Edge detection by subtraction



smoothed (5x5 Gaussian)

Edge detection by subtraction




smoothed – original
(scaled by 4, offset +128)

Why does
this work?

filter demo

Gaussian - image filter



The Canny edge detector

This is probably the most widely used edge detector in computer vision



original image (Lena)

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

The Canny edge detector



norm of the gradient

The Canny edge detector




thresholding

Get Orientation at Each Pixel



thresholding

Get Orientation at Each Pixel



$$\text{theta} = \text{atan2}(-g_y, g_x)$$

thresholding



The Canny edge detector



thinning

(non-maximum suppression)


Non-maximum suppression



Check if pixel is local maximum along gradient direction

- requires checking interpolated pixels p and r

Non-maximum suppression



Check if pixel is local maximum along gradient direction

- requires checking interpolated pixels p and r

Canny Edges




Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$



Canny with $\sigma = 2$


Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$




Canny with $\sigma = 2$


The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

An edge is not a line...



An edge is not a line...



How can we detect **lines**?


Finding lines in an image

- Option 1:
 - Search for the line at every possible position/orientation
 - What is the cost of this operation?

Finding lines in an image


- Option 1:
 - Search for the line at every possible position/orientation
 - What is the cost of this operation?
- Option 2:
 - Use a voting scheme: hough transform

Finding lines in an image




- Connection between image (x,y) and Hough (m,b) spaces
 - A line in the image corresponds to a point in Hough space

Finding lines in an image




- Connection between image (x,y) and Hough (m,b) spaces
 - A line in the image corresponds to a point in Hough space
 - To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$

Finding lines in an image



- Connection between image (x,y) and Hough (m,b) spaces
 - A line in the image corresponds to a point in Hough space
 - To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$
 - What does a point (x_0, y_0) in the image space map to?

Finding lines in an image



- Connection between image (x,y) and Hough (m,b) spaces
 - A line in the image corresponds to a point in Hough space
 - To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$
 - What does a point (x_0, y_0) in the image space map to?
 - A: the solutions of $b = -x_0m + y_0$
 - this is a line in Hough space

Hough transform algorithm

- Typically use a different parameterization

$$d = x \cos \theta + y \sin \theta$$

- d is the perpendicular distance from the line to the origin
- θ is the angle


Hough transform algorithm

- Basic Hough transform algorithm
 - Initialize $H[d, \theta] = 0$
 - for each edge point $I[x, y]$ in the image
 - for $\theta = 0$ to 180
 - $d = x \cos \theta + y \sin \theta$
 - $H[d, \theta] += 1$
 - Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum
 - The detected line in the image is given by $d = x \cos \theta + y \sin \theta$



Hough transform algorithm

- Basic Hough transform algorithm
 - Initialize $H[d, \theta] = 0$
 - for each edge point $I[x,y]$ in the image
 - for $\theta = 0$ to 180
 - $d = x\cos\theta + y\sin\theta$
 - $H[d,\theta] += 1$
 - Find the value(s) of (d,θ) where $H[d,\theta]$ is maximum
 - The detected line in the image is given by $d = x\cos\theta + y\sin\theta$
- What's the running time (measured in #votes)?



Hough transform algorithm





Hough transform algorithm



Hough transform algorithm



Hough transform algorithm



Extensions

- Extension 1: Use the image gradient
 - same
 - for each edge point $I[x,y]$ in the image
 - compute unique (d, θ) based on image gradient at (x,y)
 - $H[d, \theta] += 1$
 - same
 - same
- What's the running time measured in votes?
- Extension 2
 - give more votes for stronger edges
- Extension 3
 - change the sampling of (d, θ) to give more/less resolution
- Extension 4
 - the same procedure can be used with circles, squares, or any other shape, How?

A large, solid blue abstract shape that spans across the upper half of the slide. It has a jagged, angular top edge and a more regular bottom edge, creating a stylized, modern graphic element.

CONCORDIA.CA